

NAMAN GUPTA

+91 8630137824 | nmn.gupta.0515@gmail.com | github.com/itz-naman007 | linkedin.com/in/gupta-naman69
leetcode.com/u/pZDE2amUpZ/ | codolio.com/profile/PtejZltn | kaggle.com/nmngup | gupta-naman.vercel.app/

AI Engineer | Python, PyTorch, LangChain, RAG, NLP, LLM

PROFILE SUMMARY

AI/ML Engineering student building end-to-end deep learning and NLP systems, with hands-on experience in Retrieval Augmented Generation (RAG) pipelines using LangChain, FAISS, and Hugging Face Transformers. Built a RAG-based document Q&A system in Flask and a natural-language-to-SQL chatbot using LangGraph and MySQL. Comfortable across the PyTorch, Scikit-Learn, and Docker stack for building and deploying ML systems.

EDUCATION

B.Tech, Computer Science (AI Specialization) Bennett University, Greater Noida | Expected July 2027 | CGPA: 7.27

KEY SKILLS

Programming & Databases: Python, C++, MySQL, MongoDB

ML & Deep Learning: PyTorch, TensorFlow, Scikit-Learn, CNNs, RNNs (LSTM, GRU, BiGRU), Computer Vision, GANs, Fine-tuning

LLM & NLP: LangChain, LangGraph, Hugging Face Transformers, RAG, Agentic AI, Prompt Engineering, NLP, NLTK, GloVe

Vector Databases: FAISS, ChromaDB, Postgresql, Redis

Deployment & Tools: Flask, FastAPI, Django, Docker, Git, CI/CD, REST APIs, MLFlow

PROJECTS

Chat-With-Documents: RAG Pipeline

Python • Flask • FAISS • HuggingFace [GitHub](#) | Mar 2026

- Built a production-grade RAG pipeline with FAISS and a custom-made pipeline, enabling multi-lingual semantic Q&A over documents; integrated Hugging Face transformer embeddings and conversation memory to support coherent multi-turn dialogue.
- Engineered a full document preprocessing system (chunking, embedding, indexing for 5+ formats) and a Flask backend with query routing across three modes Summarize, Q&A reducing average query latency.

Fake News Detection: NLP Deep Learning

Python • TensorFlow • NLTK • Scikit-Learn [GitHub](#) | Dec 2025

- Designed and benchmarked 4 deep learning classifiers (GRU, BiGRU, LSTM, CNN+LSTM) on the WELFake dataset (~72K articles) with full NLP preprocessing — tokenisation, GloVe embeddings, sequence padding — achieving 98%+ accuracy.
- Conducted a systematic ablation study across architectures and hyperparameters, documenting precision/recall trade-offs; BiGRU delivered the highest F1 score with the fewest false positives.

Text-to-SQL Chatbot

Python • LangChain • MySQL • Gradio [GitHub](#) | Sep 2026

- Built a natural-language-to-SQL chatbot using Google Gemini and LangGraph to convert user queries into SQL and retrieve results from a MySQL database.
- Integrated the LangChain SQL Toolkit and Gradio to enable database schema inspection, dynamic SQL execution, and an interactive web-based interface.

CERTIFICATIONS & ACHIEVEMENTS

- Machine Learning for All (University of London) [Certificate](#)
- PyTorch: Fundamentals (DeepLearning.AI) [Certificate](#)
- Retrieval Augmented Generation (RAG) (DeepLearning.AI) [Certificate](#)
- Database Management Essentials (University of Colorado System) [Certificate](#)
- Solved 150+ questions on LeetCode and other coding platforms